



方圆并济：基于 Spark on Angel 的高性能分布式机器学习

Tencent——数据平台部

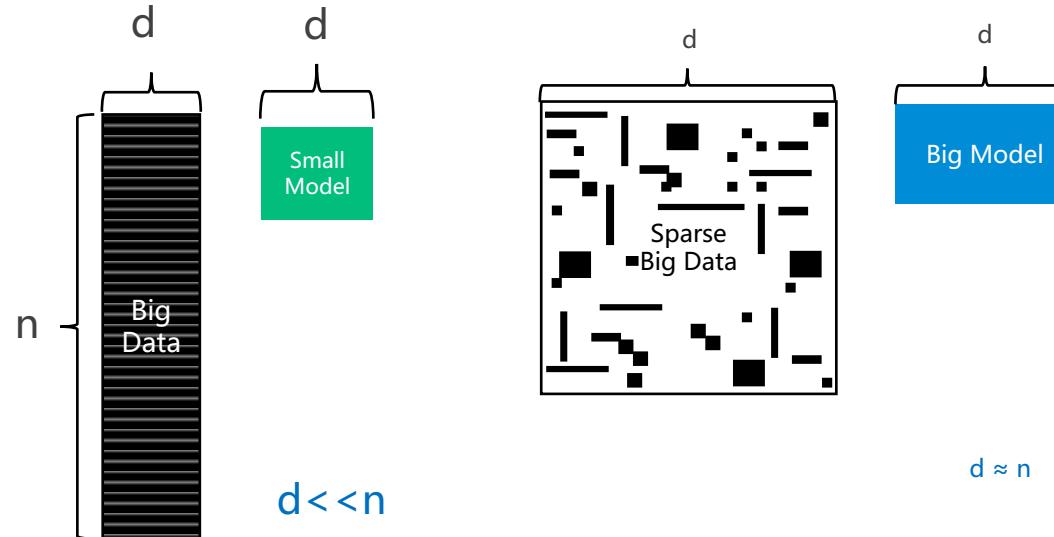
Andymhuang (黄明)

目录

- 源起
- Spark on Angel
- Spark on Angel的开发
- Spark on Angel的算法
- 性能和比较
- 后续规划

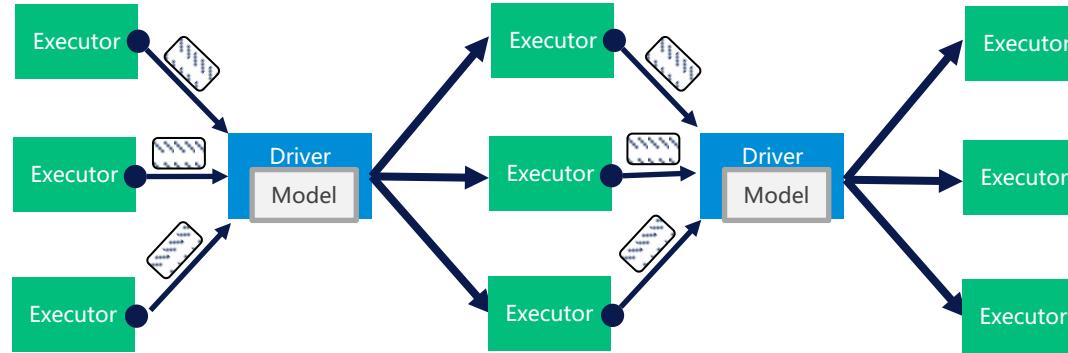
源起

腾讯的产品需求



寻找满足十亿级维度的工业级的分布式机器学习平台

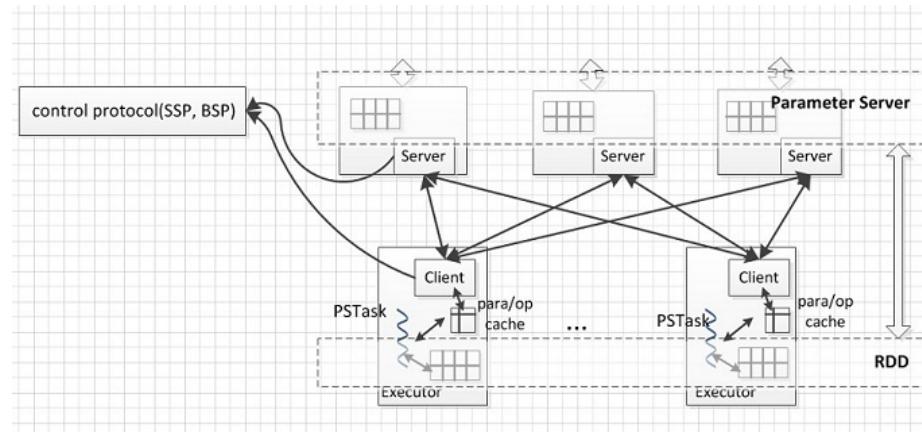
Spark机器学习的瓶颈



- Driver成为参数汇总的单点瓶颈，难以支撑大规模模型及数据
- 满足不了十亿级维度的模型训练，实际应用中需要进行降维处理
- Executor之间相互等待，整体效率不高

One Issue

A Prototype of Parameter Server



<https://issues.apache.org/jira/browse/SPARK-6932>

2015

Glint & Yahoo



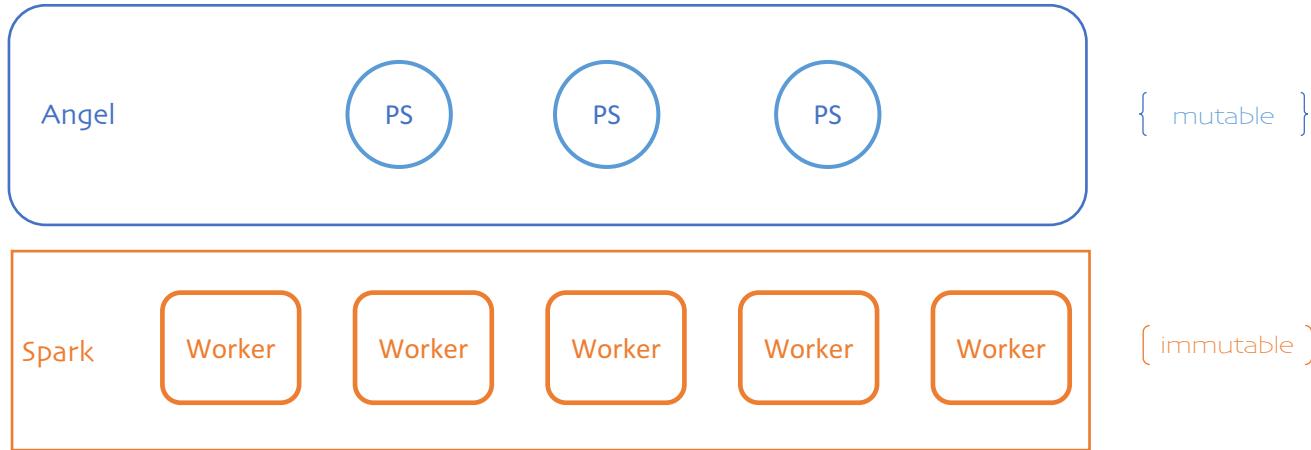
Scaling Machine Learning To Billions of Parameters

Badri Bhaskar, Erik Ordentlich
(joint with Andy Feng, Lee Yang, Peter Cnudde)
Yahoo, Inc.



2016

理念 — 方圆并济



1. 分离系统中的变和不变
2. 以少博多
3. 对Spark Core的侵入性越少越好

Spark on Angel

核心抽象

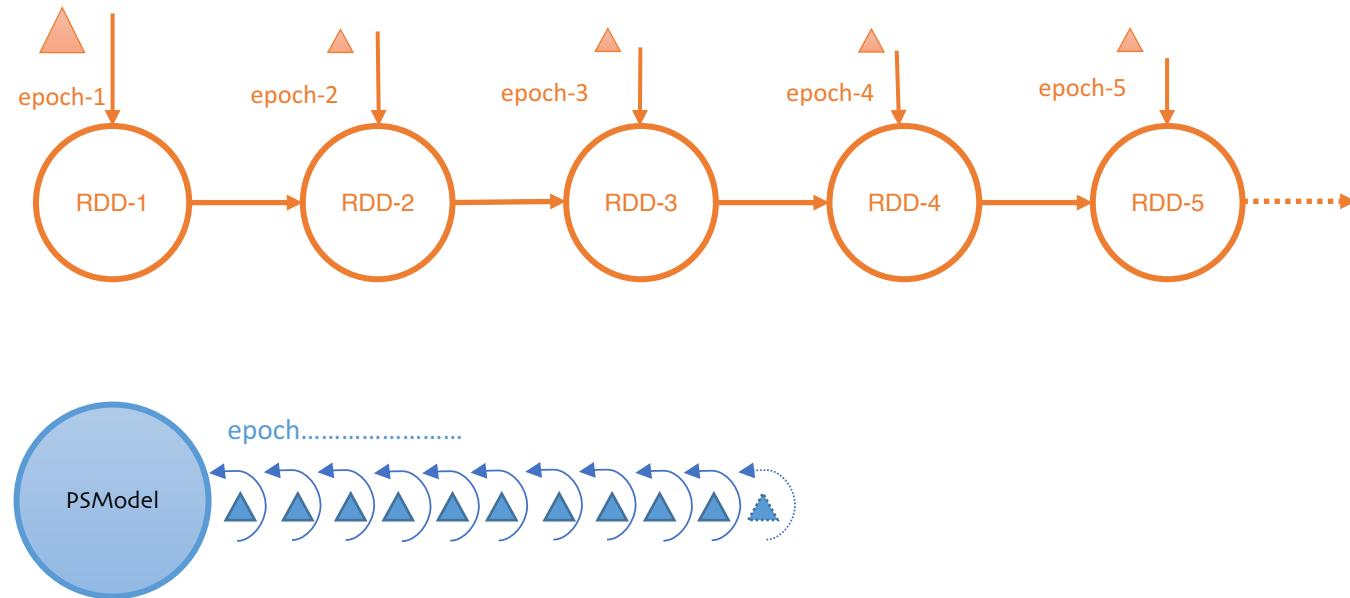
Mapper
Reducer

RDD

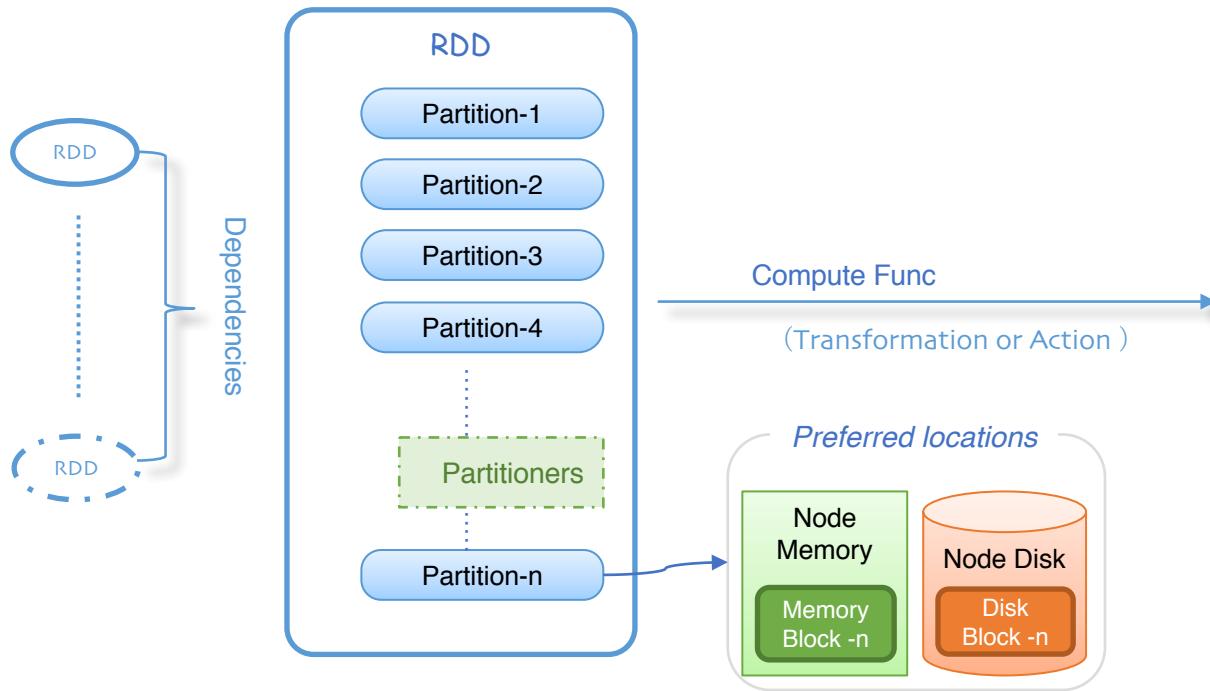
PSModel



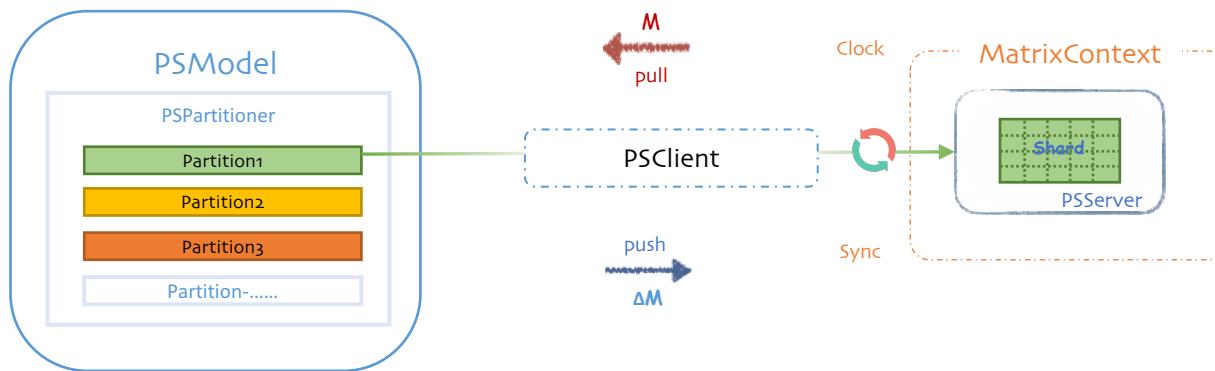
RDD vs PSModel



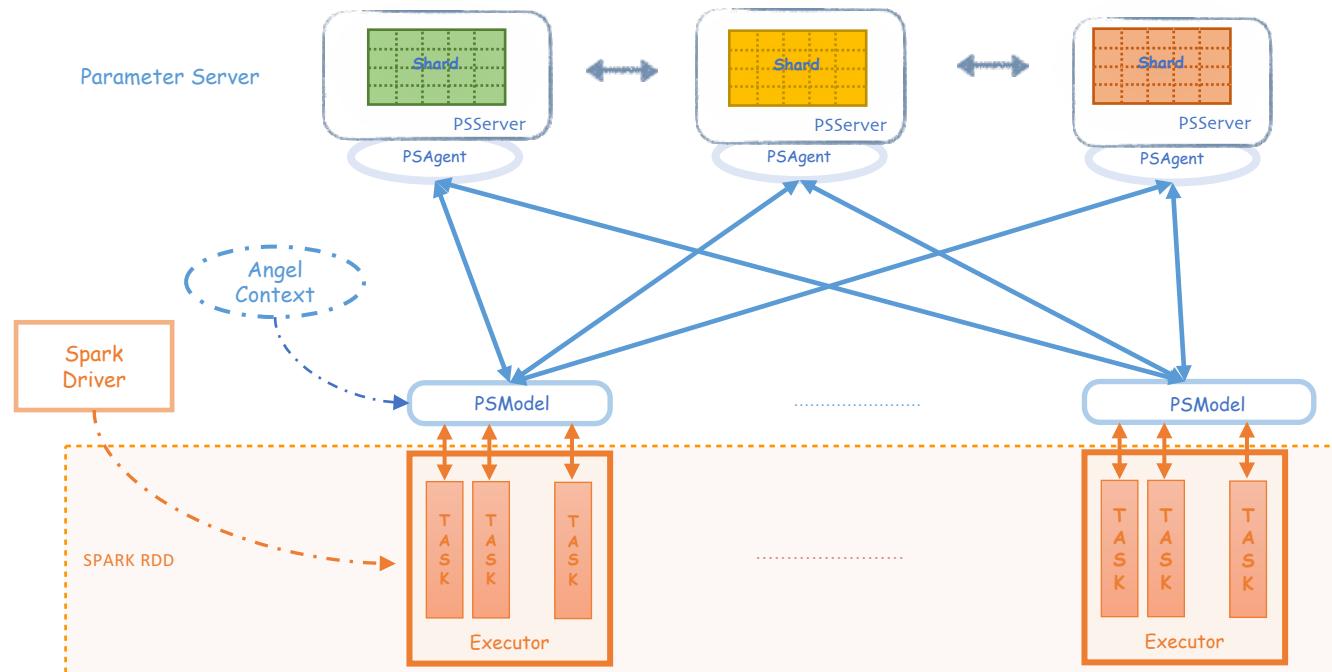
RDD的核心抽象



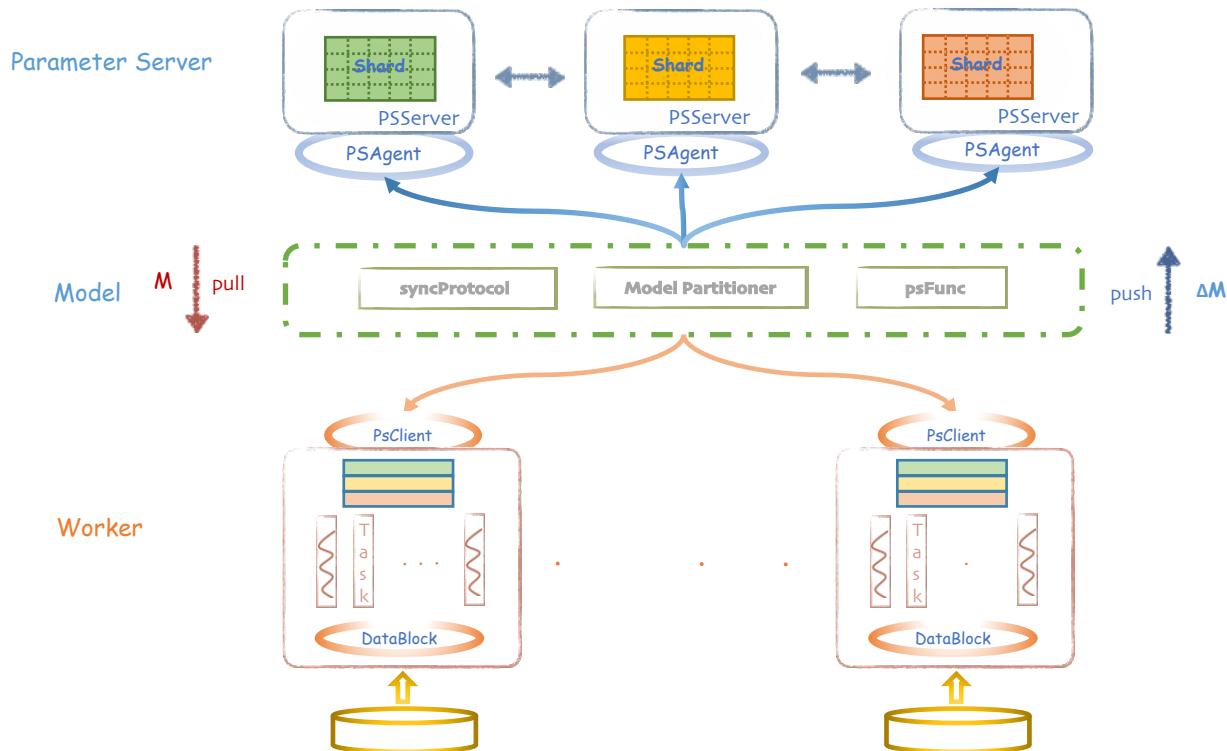
PSModel的核心抽象



Spark on Angel的架构



Angel的系统框架





开始研发

• 2015

正式开源 V1.0.0

• 2017

投入生产

• 2016

- 能支持十亿级别维度的模型训练
- 基于Matrix/Vector的模型自动切分和管理，兼顾稀疏和稠密两种格式
- 提供多种同步控制机制（BSP/SSP/ASP）

工业级别可用的
参数服务器

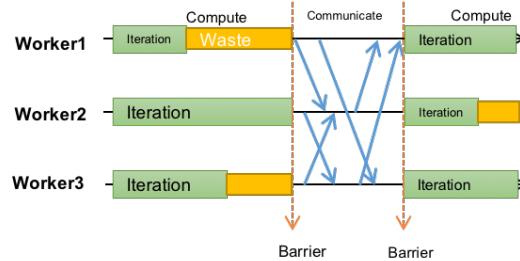
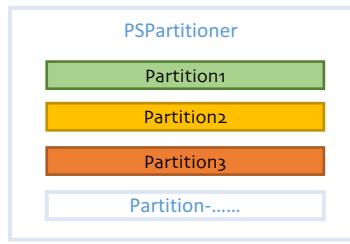
丰富的机器学习及
数学计算库

- 集成LR (ADMM-LR) , SVM, KMeans, LDA, MF, GBDT等机器学习算法
- 多种优化方法，包括ADMM, OWLQN, LBFGS和GD
- 支持多种损失函数、评估指标，包含L1、L2正则项算法

- 基于PSModel的机器学习友好接口，以Model为核心编程
- 支持Spark on Angel，Spark代码少量改动就可以运行Angel之上
- 灵活的psFunc，便于复杂算法的开发，实现模型并行

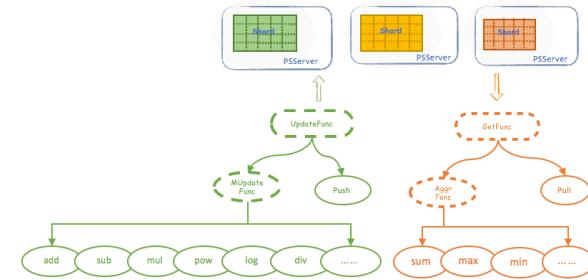
友好的
用户编程接口

Angel和Glint的比较



更丰富的模型切分

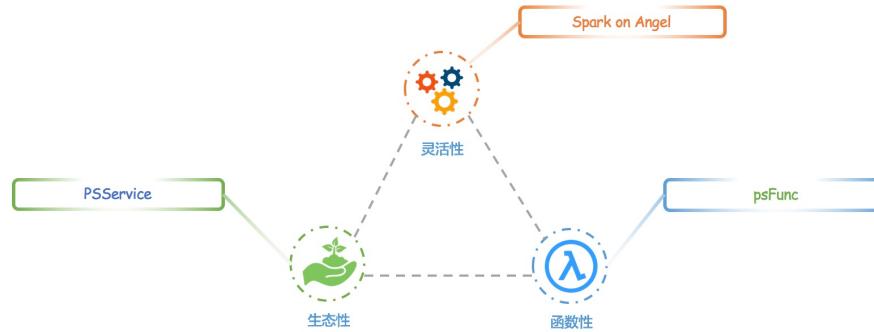
更灵活的异步模式



更强大的psFunc

Angel的定位

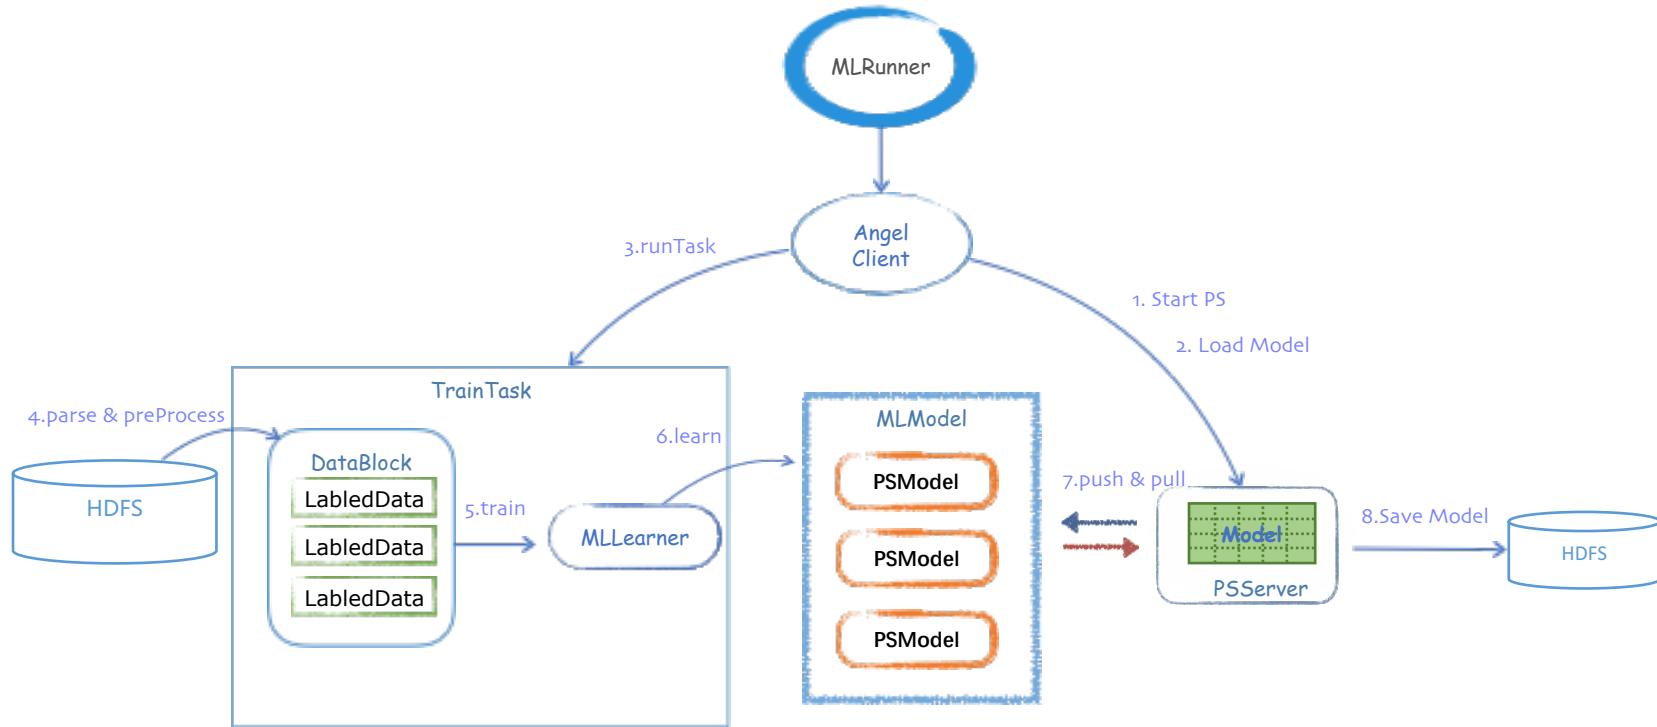
一个源于Parameter Server理念的高性能分布式机器学习平台



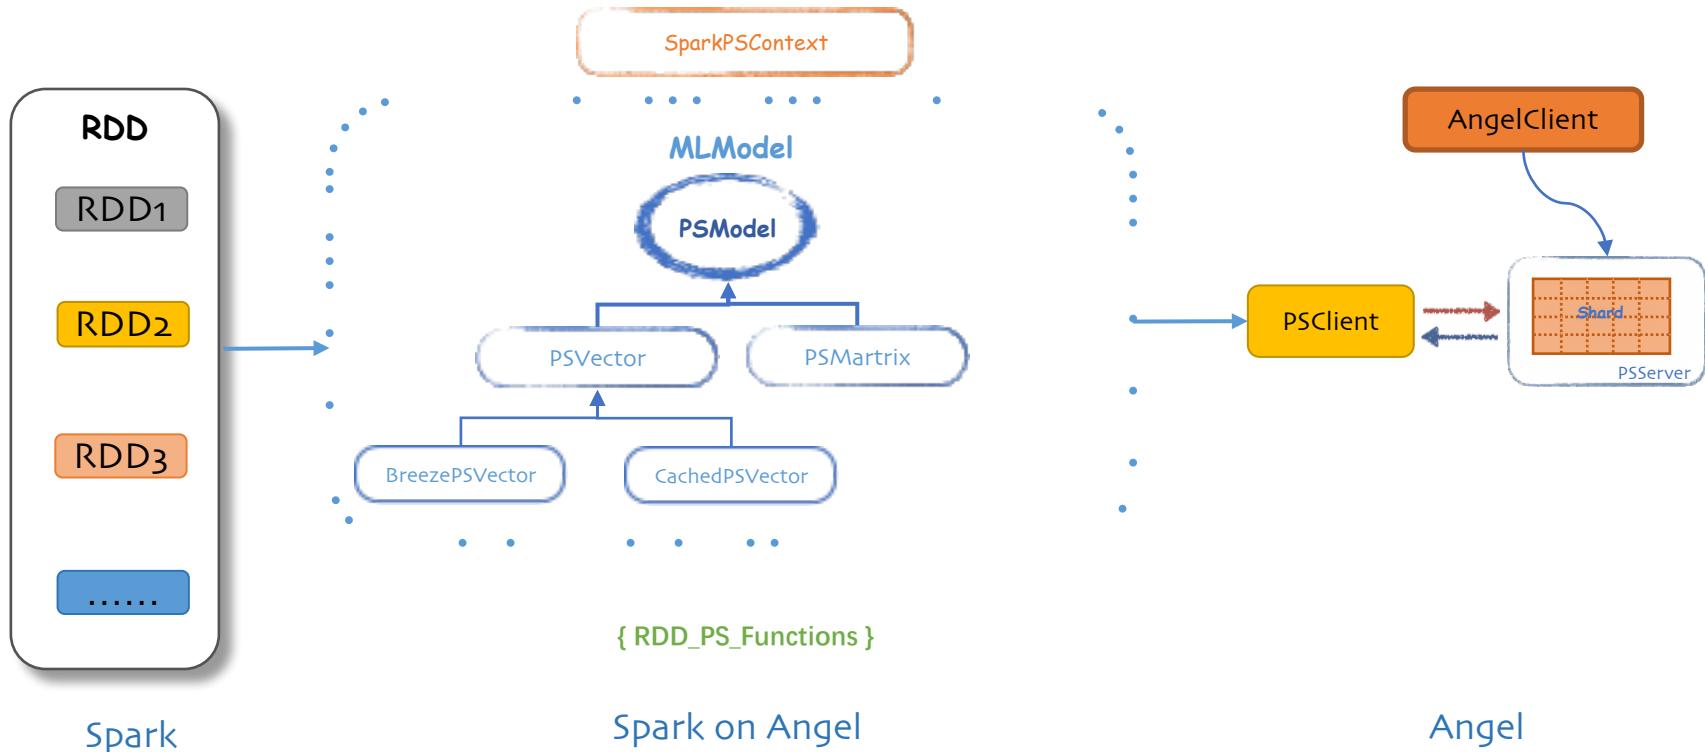
<https://github.com/tencent/angel>

Spark on Angel的开发

Angel的API设计



Spark on Angel的API设计



Spark on Angel的基础写法

```
val psContext = PSContext.getOrCreate(spark.sparkContext)
val psVector = PSVector.dense(0.0)
rdd.map { case (label, feature) =>
    psVector.increment(feature)
    ...
}
println("Feature sum size:" + psVector.dimension)
```

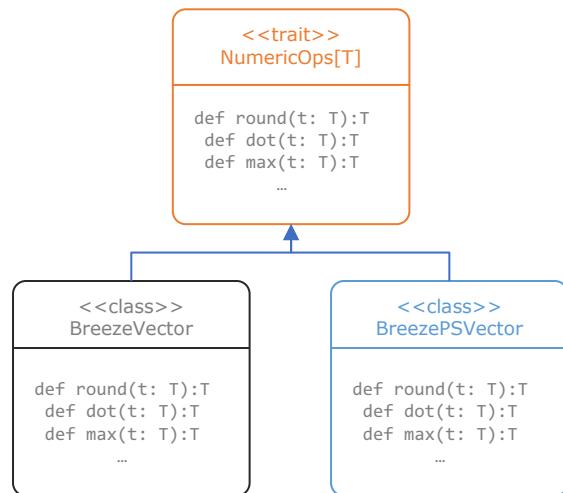
- 启动SparkSession

- 初始化PSContext，启动Angel的PSServer
- 通过PSContext, 创建PSVector
- 在RDD的运算中，直接调用PSVector，进行模型更新
- 终止PSContext

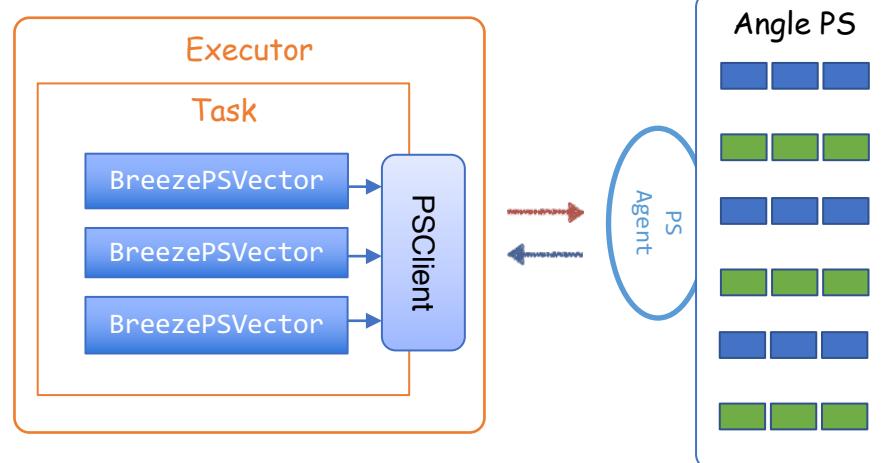
- 停止SparkSession

Vector的透明替换

混入相同特征



进行透明替换



- 将BreezeVector透明替换为BreezePSVector
- 适用于MLLib大部分算法
- 替代成本非常低

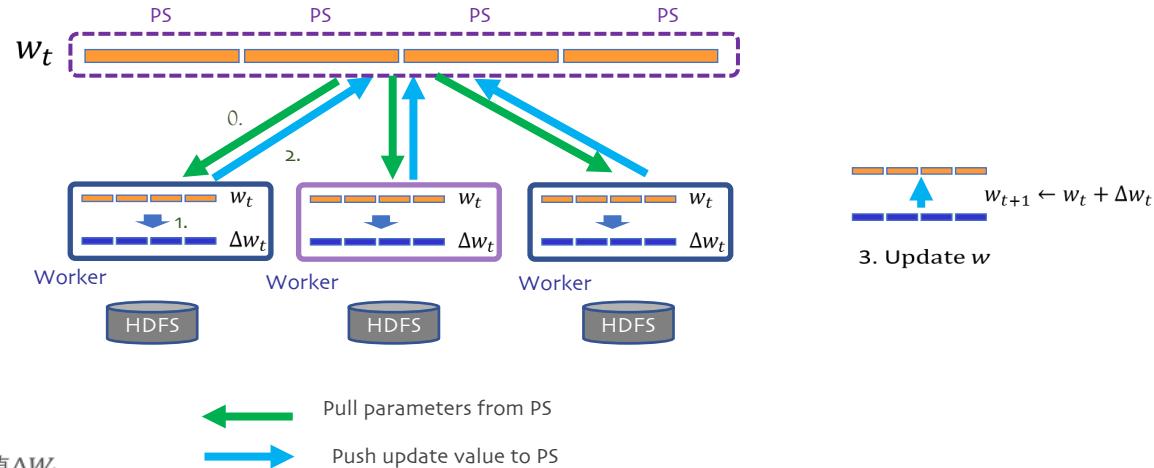
Spark on Angel的算法

Angel的算法



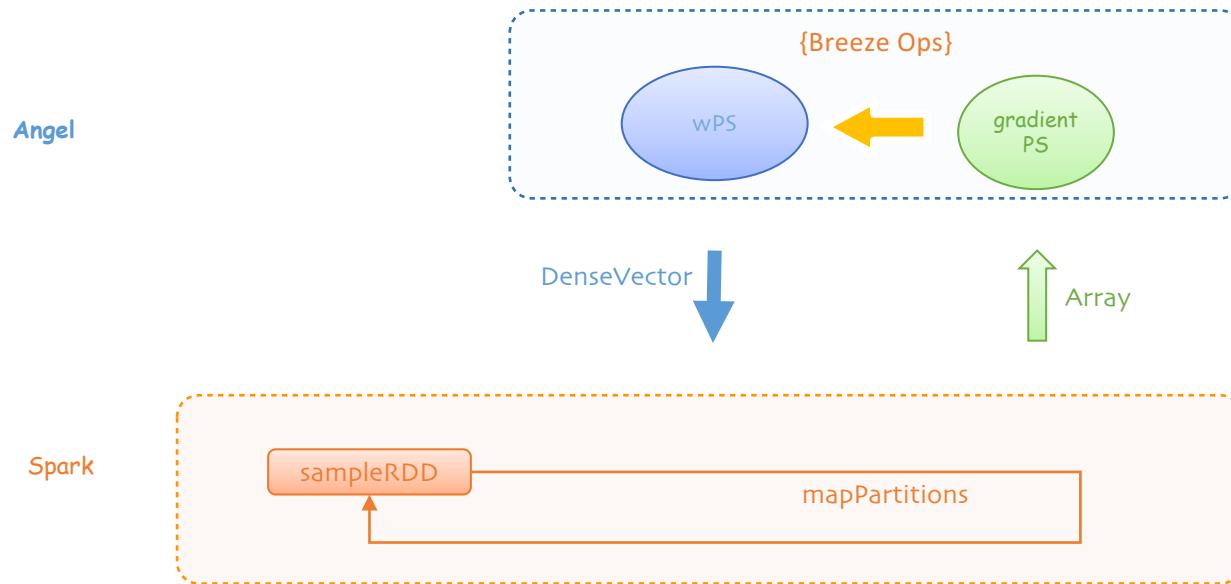
○ Spark on Angel Available

LR on Angel



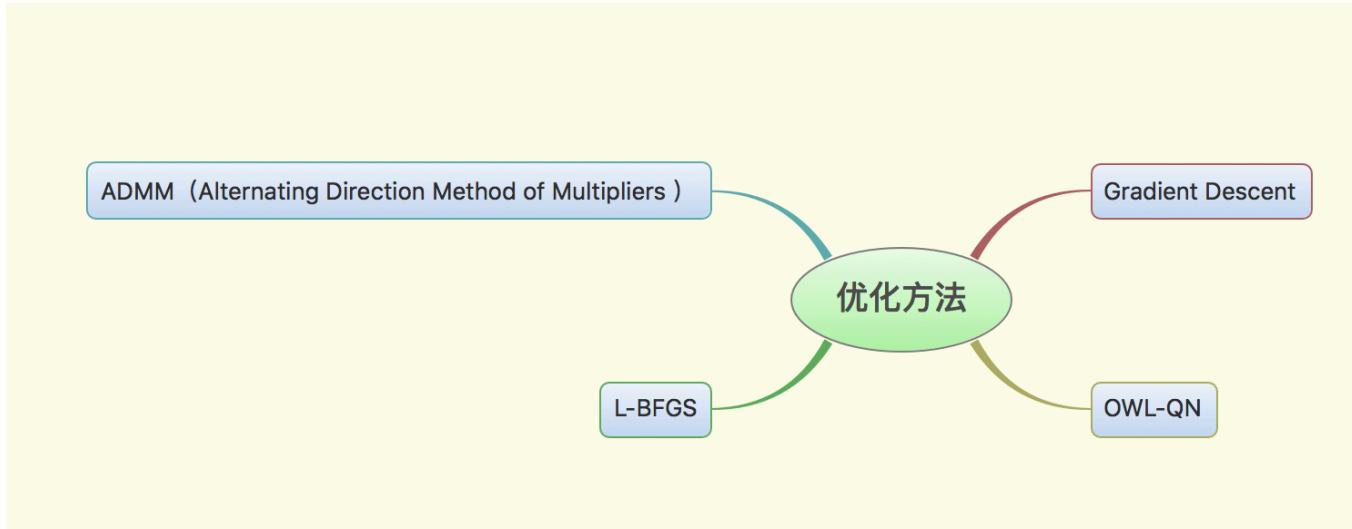
- Step 0: Worker从PS获得参数 W_t
- Step 1: Worker计算参数的更新值 ΔW_t
- Step 2: Worker把 ΔW_t 推送给PS
- Step 3: PS更新参数 ($W_{t+1} \leftarrow W_t + \Delta W_t$)

[Spark on Angel] LR

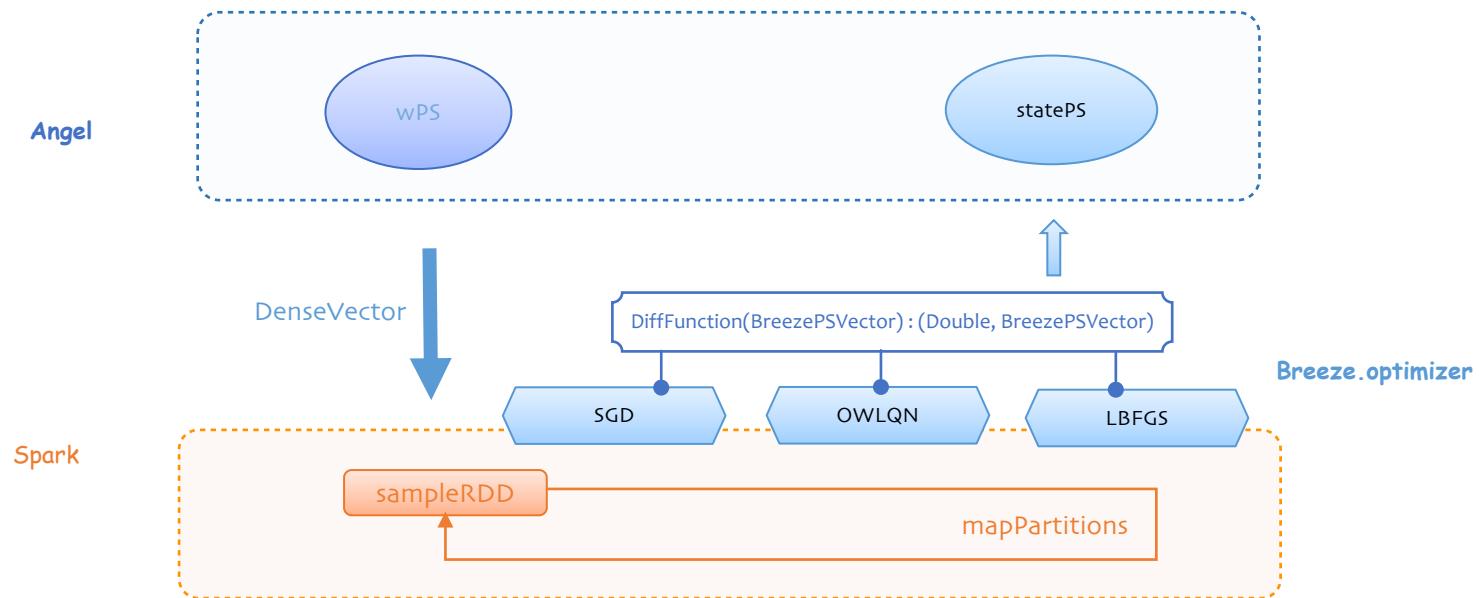


[spark_on_angel_quick_start.md]

优化方法



[Spark on Angel] LR with Optimizer

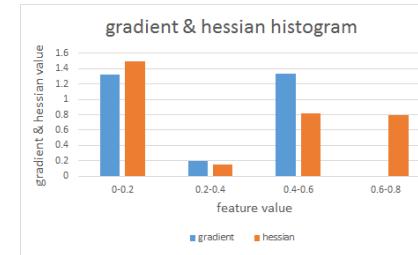


[spark_on_angel_optimizer.md]

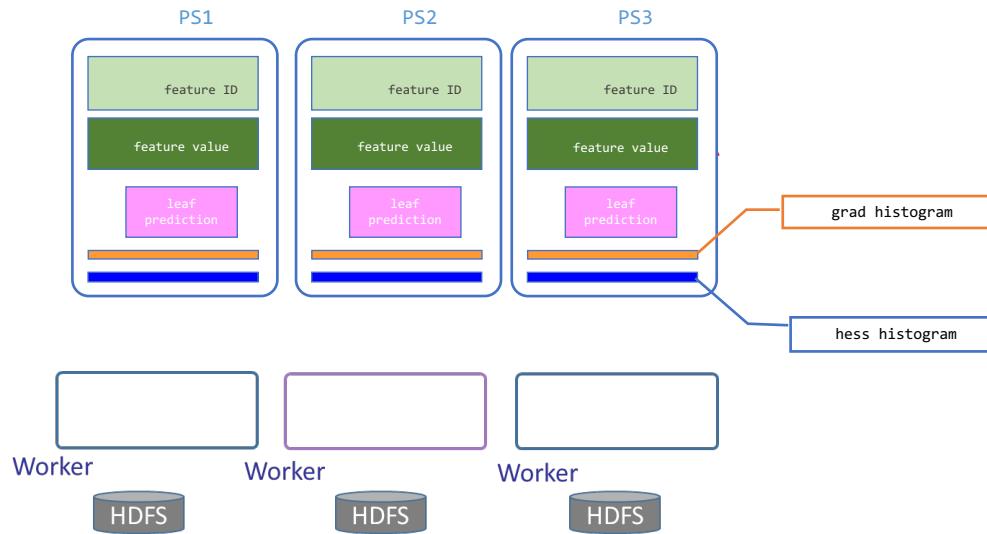
GBDT : 树模型+Boosting



- A predict() $5 + 0.5 = 5.5$
- B predict() $10 + 1.5 = 11.5$
- C predict() $1 + 1.5 = 2.5$
- D predict() $1 + 0.5 = 1.5$
- E predict() $1 + 1.5 = 2.5$

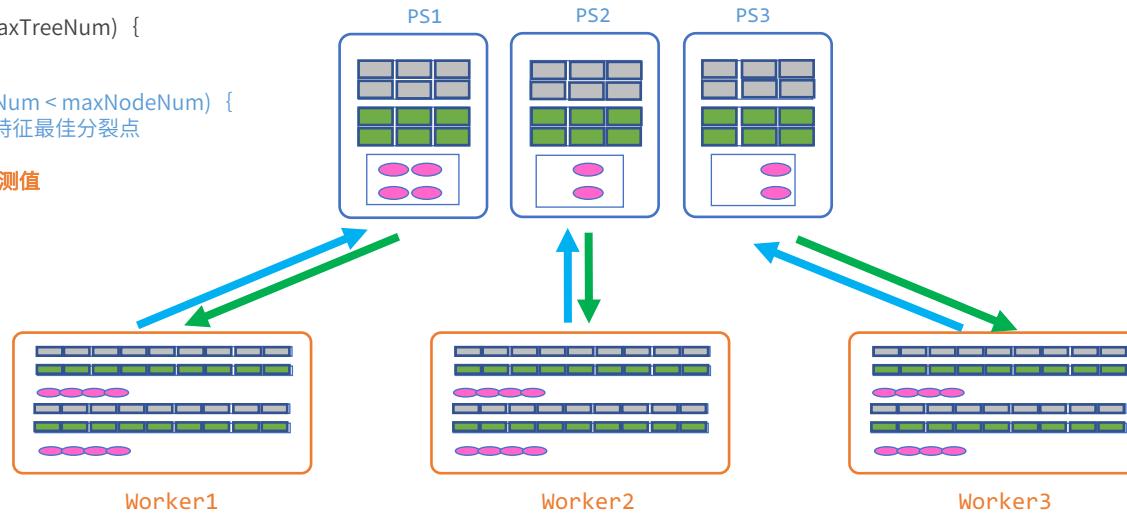


GBDT on Angel: 模型存储



GBDT on Angel (1) : 构建森林

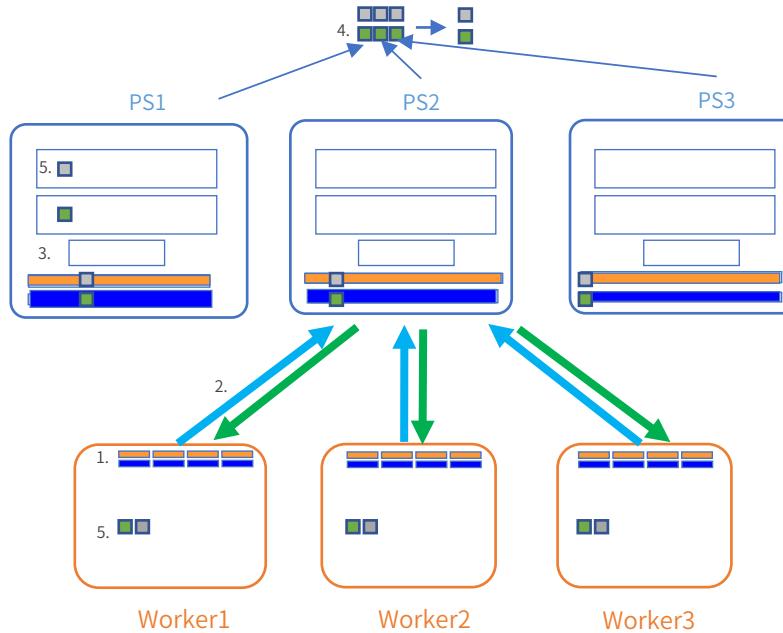
```
while (treeNum < maxTreeNum) {  
    创建一棵新树  
    while (nodeNum < maxNodeNum) {  
        寻找特征最佳分裂点  
    }  
    计算叶子节点的预测值  
    完成一棵决策树  
}
```



GBDT on Angel (2) : 分裂树节点

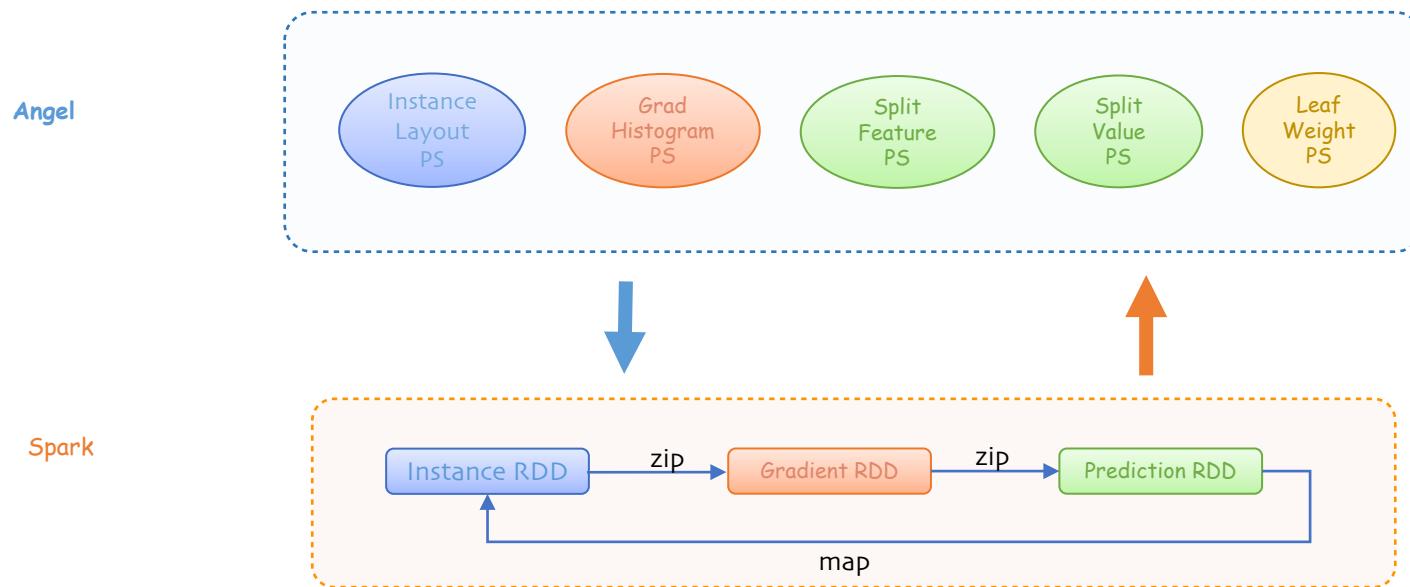
find split feature & value

1. worker计算梯度直方图（一阶&二阶）
2. worker推送梯度直方图到PS
3. 每个PS计算局部最佳分裂点
4. PS之间计算出全局最佳分裂点
5. 创建分裂点，Worker从Ps拉取最佳分裂点



[gbdt_on_angel.md]

[Spark on Angel] GBDT



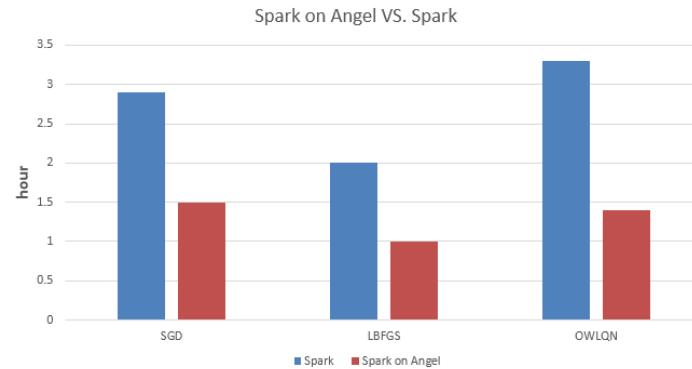
[spark_on_angel_gbdt.md]

性能比对

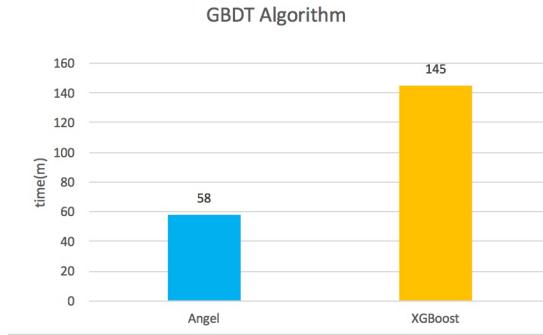
——生产数据，现网环境，尽量公平

(Spark on Angel) vs Spark —— LR

	Spark	Spark on Angel	加速比例
SGD LR (stepSize=0.05,maxIter=100)	2.9 hour	1.5 hour	48.3%
L-BFGS LR (m=10, maxIter=50)	2 hour	1 hour	50.0%
OWL-QN LR (m=10, maxIter=50)	3.3 hour	1.4 hour	57.6%



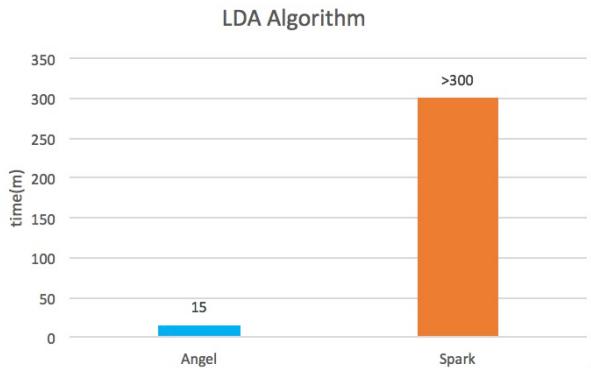
Angel vs XGBoost —— GBDT



框架	Worker	PS	建立20棵树时间
Angel	50 个(内存: 10G / Worker)	10个 (内存: 10G / PS)	58 min
XGBoost	50个 (内存: 10G / Worker)	N/A	2h 25 min

数据：腾讯内部某性别预测数据集， 3.3×10^5 特征， 1.2×10^8 样本

Angel vs Spark —— LDA



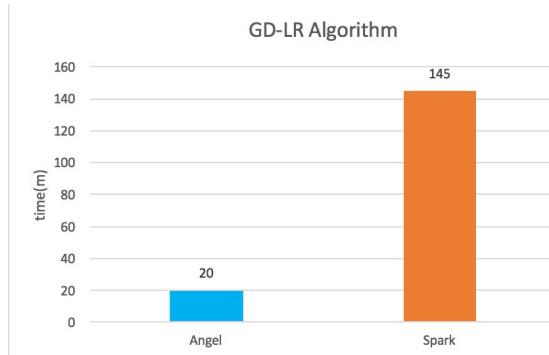
框架	Worker	PS	时间
Angel	20个(内存: 8G/Worker)	20个(内存: 4G/PS)	15min
Spark	20个(内存: 20G/Worker)	N/A	>300min

数据: PubMED

框架	Worker	PS	时间
Angel	50个(内存: 10G/Worker)	50个(内存: 4G/PS)	1h 7min

DataSet: 40G Token: 2 billion
Word: 52w Topic : 1000

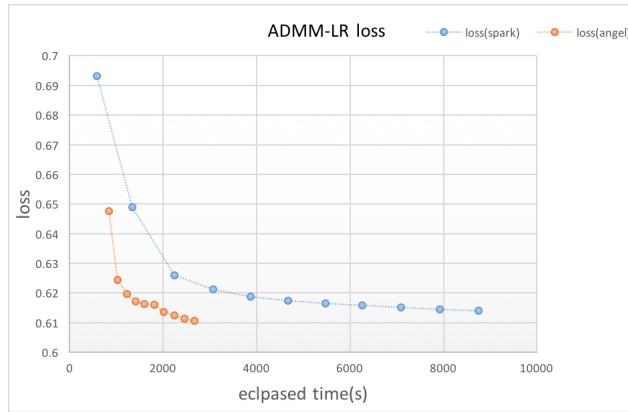
Angel vs Spark —— GD-LR



框架	Worker	PS	迭代100次时间
Angel	50个(内存:10G/Worker)	20个(内存: 5G/PS)	20min
Spark	50个(内存:14G/Worker)	N/A	145min

数据：腾讯内部某推荐数据， 5×10^7 特征， 8×10^7 样本

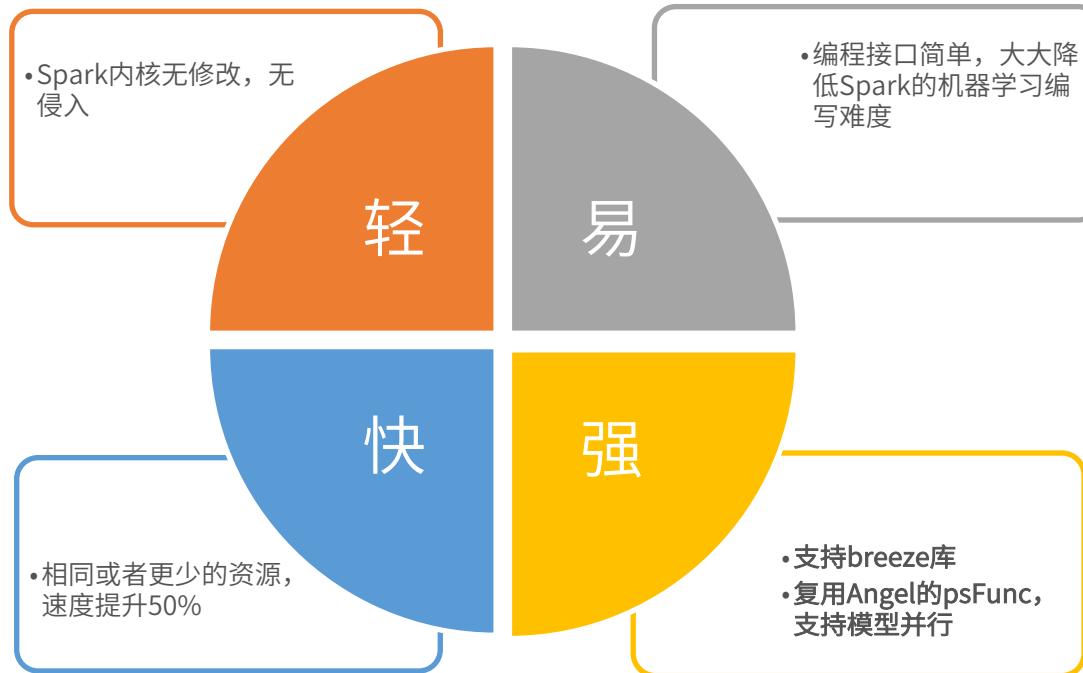
Angel vs Spark —— ADMM-LR



框架	Worker	PS	收敛退出
Angel	100个(内存:10G/Worker)	50个(内存: 5G/PS)	27 min
Spark	200个(内存:20G/Worker)	N/A	145 min

数据：腾讯内部某推荐数据，5千万特征，1亿样本

Spark on Angel的特点



开源和展望

OpenSource & Perspective

Angel开源

Tencent / angel

Unwatch ▾

289

Unstar

2,582

Fork

613

github issues

(PR 60)

- [LightBGM作者: \[GBDT\] The purposes of using parameter server in GBDT #7](#)
- [海外华人: English translation of documents #95](#)
- [华为工程师: \[WIP\]Upgrade the netty version of RPC to 4.x #94](#)
- [新浪微博: 增强LR算法, 加入y截距因子](#)
-

学术创新

- 国际顶级会议Paper (CCF A类)

- [LDA*: A Robust and Large-scale Topic Modeling System VLDB, 2017](#)
- [Heterogeneity-aware Distributed Parameter Servers. SIGMOD, 2017](#)
- Angel: a new large-scale machine learning system. National Science Review (NSR), 2017
- TencentBoost: A Gradient Boosting Tree System with Parameter Server. ICDE, 2017
-



版本展望 (What is Next)



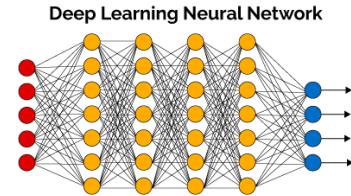
Python API

V1.3



Spark Streaming on Angel

V1.5



Deep Learning Framework Support

V2.0

Q & A

微博: @明风

喜欢记得给个Star噢



andymhua@tenant.com

机器学习系统 & 算法工程师

We are Hiring